

Theano at a Glance: A Framework for Machine Learning

Frédéric Bastien

Montreal Institute for Learning Algorithms
Université de Montréal
Montréal, Canada
`bastienf@iro.umontreal.ca`

Presentation prepared with Pascal Lamblin and Simon Lefrançois



AI With The Best 2016

Université 
de Montréal

Introduction

Introduction

Community

Theano

Description

Compiling/Running

Modifying expressions

GPU

Debugging

Models

Logistic Regression

Convolution

Scan

News

End

High level

Python <- {NumPy/SciPy/libgpuarray} <- Theano <- {...}

- ▶ Python: OO coding language
- ▶ Numpy: n -dimensional array object and scientific computing toolbox
- ▶ SciPy: sparse matrix objects and more scientific computing functionality
- ▶ libgpuarray: GPU n -dimensional array object in C for CUDA and OpenCL
- ▶ Theano: compiler/symbolic graph manipulation
 - ▶ **(Not specific to machine learning)**
- ▶ {...}: Many libraries built on top of Theano

What Theano provides

- ▶ Lazy evaluation for performance
- ▶ GPU support
- ▶ Symbolic differentiation
- ▶ Automatic speed and stability optimization

High level

Many [machine learning] library build on top of Theano

- ▶ Keras
- ▶ blocks
- ▶ lasagne
- ▶ sklearn-theano
- ▶ PyMC 3
- ▶ theano-rnn
- ▶ Morb
- ▶ ...

Goal of the stack

Fast to develop
Fast to run



Some models build with Theano

Some models that have been build with Theano.

- ▶ Neural Networks
- ▶ Convolutional Neural Networks, AlexNet, OverFeat, GoogLeNet
- ▶ RNN, CTC, LSTM, GRU
- ▶ NADE, RNADE, MADE
- ▶ Autoencoders
- ▶ Generative Adversarial Nets
- ▶ SVMs
- ▶ **many variations of above models and more**

Project status

- ▶ Mature: Theano has been developed and used since January 2008 (8 yrs old)
- ▶ Driven hundreds of research papers
- ▶ Good user documentation
- ▶ Active mailing list with worldwide participants
- ▶ Core technology for Silicon-Valley start-ups
- ▶ Many contributors (some from outside our institute)
- ▶ Used to teach many university classes
- ▶ Has been used for research at big companies
- ▶ Theano 0.8.2 released 21th of April, 2016

Theano: deeplearning.net/software/theano/

Deep Learning Tutorials: deeplearning.net/tutorial/

Theano community

Active community

- ▶ Many people reply on our mailing lists
- ▶ Hundreds of answered questions on StackOverflow
- ▶ 141 contributors to Theano 0.8
- ▶ Main developers at MILA

MILA

Institut des algorithmes d'apprentissage de Montréal
Montreal Institute for Machine Learning

- ▶ Professors 7
- ▶ Staff 7
- ▶ Visiting Scientists 1
- ▶ Post-Doc 6
- ▶ Ph.D. students 42
- ▶ M.Sc. students 22
- ▶ Interns 22
- ▶ Total 107

mila.umontreal.ca

formely known as LISA



MILA Partners

- ▶ Many universities
- ▶ Ubisoft
- ▶ Nuance
- ▶ IBM
- ▶ Google
- ▶ Facebook
- ▶ NVIDIA
- ▶ Huawei
- ▶ Intel
- ▶ D-Wave
- ▶ ApSTAT
- ▶ Qualcomm
- ▶ Imagia
- ▶ Sulfur Heron

Python

- ▶ General-purpose high-level OO interpreted language
- ▶ Emphasizes code readability
- ▶ Comprehensive standard library
- ▶ Dynamic type and memory management
- ▶ Easily extensible with C
- ▶ Slow execution
- ▶ Popular in *web development* and *scientific communities*

Introduction

Introduction

Community

Theano

Description

Compiling/Running

Modifying expressions

GPU

Debugging

Models

Logistic Regression

Convolution

Scan

News

End

Description

High-level domain-specific language for numeric computation.

- ▶ **Syntax as close to NumPy as possible**
- ▶ **Compiles** most common expressions **to C for CPU and/or GPU**
- ▶ **Limited expressivity** means more opportunities for optimizations
 - ▶ **Strongly typed** -> compiles to C
 - ▶ **Array oriented** -> easy parallelism
 - ▶ Support for **looping and branching** in expressions
 - ▶ No subroutines -> **global optimization**
- ▶ **Automatic speed and numerical stability optimizations**

Description (2)

- ▶ **Symbolic differentiation and R op** (Hessian Free Optimization)
- ▶ Can **reuse other technologies** for best performance
 - ▶ CUDA, CuBLAS, CuDNN, BLAS, SciPy, PyCUDA, Cython, Numba, ...
- ▶ Sparse matrices (CPU only)
- ▶ Extensive unit-testing and self-verification
- ▶ Extensible (You can create new operations as needed)
- ▶ Works on **Linux, OS X and Windows**
- ▶ **Multi-GPU** (via platoon)
- ▶ New GPU back-end:
 - ▶ **Float16** new back-end (need cuda 7.5)
 - ▶ **Multi dtypes**
 - ▶ Multi-GPU support in the same process

Simple example



```
import theano
# declare symbolic variable
a = theano.tensor.vector("a")

# build symbolic expression
b = a + a ** 10

# compile function
f = theano.function([a], b)

# Execute with numerical value
print f([0, 1, 2])
# prints 'array([0, 2, 1026])'
```


Simple example



Overview of library

Theano is many things

- ▶ Language
- ▶ Compiler
- ▶ Python library

Compiling and running expression

- ▶ `theano.function`
- ▶ shared variables and updates
- ▶ compilation modes

Shared variable example

```
>>> from theano import shared
>>> x = shared(0.) # Normally, model parameters
>>> updates = [(x, x + 1)]
>>> f = function([], updates=updates)
>>> f()
>>> x.get_value()
1.0
>>> x.set_value(100.)
>>> f()
>>> x.get_value()
101.0
```

Modifying expressions

There are “macro” that automatically build bigger graph for you.

- ▶ `theano.grad`
- ▶ `R_op`, `L_op` for Hessian Free Optimization
- ▶ `hessian`
- ▶ `jacobian`
- ▶ clone the graph with replacement
- ▶ you can navigate the graph if you need


Those functions can get called many times, for example to get the 2nd derivative.

The grad method

```
>>> x = T.scalar('x')
>>> cost = 2. * x
>>> g = T.grad(cost, x)
```

Print the optimized graph

```
>>> f = theano.function([x], g)
>>> theano.printing.pydotprint(f)
```



Enabling GPU

- ▶ Theano's current back-end only supports 32 bit on GPU
- ▶ libgpuarray (new-backend) supports all dtype
- ▶ CUDA supports float64, but it is slow on gamer GPUs

CuDNN

- ▶ V4, V5 and V5.1 are supported.
- ▶ It is enabled automatically if available.
- ▶ Theano flag to get an error if can't be used:
“dnn.enabled=True”
- ▶ Theano flag to disable it: “dnn.enabled=False”

Debugging

- ▶ DebugMode: a mode that tests many things done by Theano (very slow)
- ▶ NanGuardMode: a mode that help find the cause of nan in the graph.
- ▶ Error message
- ▶ theano.printing.debugprint: print a textual representation of computation

Error message: code

```
import numpy as np
import theano
import theano.tensor as T
x = T.vector()
y = T.vector()
z = x + x
z = z + y
f = theano.function([x, y], z)
f(np.ones((2,)), np.ones((3,)))
```

Error message: 1st part

```
Traceback (most recent call last):
[...]
ValueError: Input dimension mismatch.
      (input [0].shape[0] = 3, input [1].shape[0] = 2)
Apply node that caused the error:
      Elemwise{add,no_inplace}(<TensorType(float64, vector)>,
                                <TensorType(float64, vector)>,
                                <TensorType(float64, vector)>)
Inputs types: [TensorType(float64, vector),
                TensorType(float64, vector),
                TensorType(float64, vector)]
Inputs shapes: [(3,), (2,), (2,)]
Inputs strides: [(8,), (8,), (8,)]
Inputs values: [array([ 1.,  1.,  1.]),
                 array([ 1.,  1.]),
                 array([ 1.,  1.])]
Outputs clients: [['output']]
```

Backtrace

Backtrace when the node is created:

File `"test.py"`, line 7, in `<module>`

```
z = z + y
```

Error message: 2nd part

HINT: Re-running with most Theano optimization disabled could give you a back-traces when this node was created. This can be done with by setting the Theano flags “optimizer=fast_compile”. If that does not work, Theano optimizations can be disabled with “optimizer=None”.

HINT: Use the Theano flag “exception_verbosity=high” for a debugprint of this apply node.

debugprint

```
>>> from theano.printing import debugprint
>>> debugprint(a)
Elemwise{mul,no_inplace} [id A]  ' '
| TensorConstant{2.0} [id B]
| Elemwise{add,no_inplace} [id C]  'z'
| <TensorType(float64, scalar)> [id D]
| <TensorType(float64, scalar)> [id E]
```

Introduction

Introduction

Community

Theano

Description

Compiling/Running

Modifying expressions

GPU

Debugging

Models

Logistic Regression

Convolution

Scan

News

End

Inputs

```
# Load from disk and put in shared variable.  
datasets = load_data(dataset)  
train_set_x, train_set_y = datasets[0]  
valid_set_x, valid_set_y = datasets[1]
```

```
# allocate symbolic variables for the data  
index = T.iscalar() # index to a [mini]batch
```

```
# generate symbolic variables for input minibatch  
x = T.matrix('x') # data, 1 row per image  
y = T.ivector('y') # labels
```


Model

```
n_in = 28 * 28  
n_out = 10
```

```
# weights
```

```
W = theano.shared(  
    numpy.zeros((n_in, n_out),  
                dtype=theano.config.floatX))
```

```
# bias
```

```
b = theano.shared(  
    numpy.zeros((n_out, ),  
                dtype=theano.config.floatX))
```

Computation

the forward pass

```
p_y_given_x = T.nnet.softmax(T.dot(input, W) + b)
```

cost we minimize: the negative log likelihood

```
l = T.log(p_y_given_x)
```

```
cost = -T.mean(l[T.arange(y.shape[0]), y])
```

the error

```
y_pred = T.argmax(p_y_given_x, axis=1)
```

```
err = T.mean(T.neq(y_pred, y))
```

Gradient and updates

```
# compute the gradient of cost
```

```
g_W, g_b = T.grad(cost=cost, wrt=(W, b))
```

```
# model parameters updates rules
```

```
updates = [(W, W - learning_rate * g_W),  
            (b, b - learning_rate * g_b)]
```

Training function

```
# compile a Theano function that train the model
train_model = theano.function(
    inputs=[index], outputs=(cost, err),
    updates=updates,
    givens={
        x: train_set_x[index * batch_size:
                        (index + 1) * batch_size],
        y: train_set_y[index * batch_size:
                        (index + 1) * batch_size]
    }
)
```

Convolution computation

```
# convolve input feature maps with filters
conv_out = conv.conv2d(input=x, filters=W)

# pool each feature map individually,
# using maxpooling
pooled_out = pool.pool_2d(
    input=conv_out,
    ds=(2, 2), # poolsize
    ignore_border=True)

output = T.tanh(pooled_out +
                 b.dimshuffle('x', 0, 'x', 'x'))
```

Scan

- ▶ Allows looping (for, map, while)
- ▶ Allows recurrence (reduce)
- ▶ Allows recurrence with dependency on many of the previous time steps
- ▶ Optimize some cases like moving computation outside of scan
- ▶ The Scan grad is done via Backpropagation Through Time (BPTT)

Theano 0.8

Released 21th of March, 2016

Highlights:

- ▶ Faster compilation and execution
- ▶ Integration of CuDNN for better GPU performance
- ▶ Many Scan improvements (execution speed up, ...)
- ▶ Interactive visualization of graphs with d3viz
- ▶ `optimizer=fast_compile` moves computation to the GPU.
- ▶ Multi-GPU for data parallism via Platoon
<https://github.com/mila-udem/platoon>
- ▶ `cnmem` (better memory management on GPU)
- ▶ New GPU back-end

A total of 141 people contributed to this release.

Since last release

Highlights:

- ▶ Computation and compilation speed up
- ▶ Multi-cores convolution and pooling on CPU
- ▶ More numerical stability by default for some graphs
- ▶ CuDNN: 5.1, batch normalization, (RNN soon)
- ▶ Dilated convolution
- ▶ Multiple-GPU, synchronous update (via platoon)
- ▶ Partial function evaluation
- ▶ Add gradient of solve, tensorinv (CPU), tensorsolve (CPU) searchsorted (CPU)
- ▶ Add Multinomial Without Replacement
- ▶ ...

90 people already contributed to the development version.

Roadmap

- ▶ Deprecate the old GPU back-end
- ▶ Faster compile time
 - ▶ First time compilation (less C code compilations, already cached)
 - ▶ Of big graph (faster graph optimization)
- ▶ Handling of bigger graph
- ▶ Speed/Memory trade off
 - ▶ Scan very long sequence, (ex: Audio recognition and generation)
 - ▶ Recompute cheap operation of the forward during the gradient
- ▶ More Multi-GPU data parallelism algo

Where to learn more

- ▶ Deep Learning Tutorials with Theano:
`deeplearning.net/tutorial`
- ▶ Theano tutorial:
`deeplearning.net/software/theano/tutorial`
- ▶ Theano website: `deeplearning.net/software/theano`
- ▶ Doc of frameworks on top of Theano like Blocks, Keras, Lasagne, ...

Questions, acknowledgments

Questions? Acknowledgments

- ▶ All people working or having worked at MILA institute (previously LISA lab).
- ▶ All Theano users/contributors.
- ▶ Compute Canada, RQCHP, NSERC, NVIDIA, and Canada Research Chairs for providing funds, access to computing resources, hardware or GPU libraries.